


Lecture 3 - Learning Basics

Dr. Xiaowei Huang

<https://cgi.csc.liv.ac.uk/~xiaowei/ai.html>

(Attendance Code: **665240**)

Research



Computer Science Review
Volume 37, August 2020, 100270



Survey

A survey of safety and trustworthiness of deep neural networks: Verification, testing, adversarial attack and defence, and interpretability ☆

Xiaowei Huang ^a✉, Daniel Kroening ^b✉, Wenjie Ruan ^c✉, James Sharp ^d✉, Youcheng Sun ^e✉, Emese Thamo ^a✉, Min Wu ^b✉, Xinping Yi ^a✉

Entering LLM/GenAI

[Home](#) > [Artificial Intelligence Review](#) > [Article](#)

A survey of safety and trustworthiness of large language models through the lens of verification and validation

[Open access](#) | Published: 17 June 2024


Volume 57, article number 175, (2024) [Cite this article](#)

[Download PDF](#) ↓

✓ You have full access to this [open access](#) article

Xiaowei Huang ✉, Wenjie Ruan, Wei Huang, Gaojie Jin, Yi Dong, Changshun Wu, Saddek Bensalem, Ronghui Mu, Yi Qi, Xingyu Zhao, Kaiwen Cai, Yanghao Zhang, Sihao Wu, Peipei Xu, Dengyu Wu, Andre Freitas & Mustafa A. Mustafa


In the last lecture,



What is machine learning?



A few applications of machine learning



consider how to represent instances as fixed-length feature vectors



Topics

- Learning basics:
 - Before learning: data collection
 - Learning tasks: supervised and unsupervised learning
 - Learning schemes

Before learning: data collection



Independent and identically distributed (i.i.d.)

- we often assume that training instances are *independent and identically distributed* (i.i.d.) – sampled independently from the same unknown distribution
- A set of random variables $X_1, X_2, \dots, X_n$ are i.i.d. if:
 - They are **independent** (no influence on each other).
 - They are **identically distributed** (drawn from the same distribution).



Independent and identically distributed (i.i.d.)


- we often assume that training instances are *independent and identically distributed* (i.i.d.) – sampled independently from the same unknown distribution
- there are also cases where this assumption does not hold
 - cases where sets of instances have dependencies
 - instances sampled from the same medical image
 - instances from time series
 - etc.

Learning tasks: supervised and
unsupervised learning

Three
canonical
learning
problems



1. Regression - supervised

- estimate parameters, e.g. of weight vs height



2. Classification - supervised


- estimate class, e.g. handwritten digit classification




Three
canonical
learning
problems

3. Unsupervised learning – model the data

- clustering





- dimensionality reduction



The supervised learning task

- problem setting
 - set of possible instances: X
 - unknown *target function*: $f : X \rightarrow Y$
 - set of *models* (a.k.a. *hypotheses*): $H = \{h \mid h : X \rightarrow Y\}$
- given *training set* of instances of unknown target function f
 $(\mathbf{x}^{(1)}, y^{(1)}), (\mathbf{x}^{(2)}, y^{(2)}) \dots (\mathbf{x}^{(m)}, y^{(m)})$
- Output
 - model $h \in H$ that best approximates target function


$$f(x) = w_1x + w_2$$



$$f(x) = (w_1x + w_2) > 10$$

The supervised learning task


$$f(x) = (w_1x + w_2) > 10$$

- when y is discrete, we term this a *classification* task (or *concept learning*)
- when y is continuous, it is a *regression* task
- there are also tasks in which each y is more structured object like a *sequence* of discrete labels (as in e.g. image segmentation, machine translation)



$$f(x) = w_1x + w_2$$



Model representations



- throughout the semester, we will consider a broad range of representations for learned models (i.e., hypothesis space), including
 - decision trees
 - neural networks
 - Linear/logistic regression
 - Bayesian networks
 - etc.



Mushroom features (from the UCI Machine Learning Repository)

Input features

Values of the features

cap-shape: bell=b,conical=c,convex=x,flat=f, knobbed=k, **sunken=s**
 cap-surface: fibrous=f,grooves=g,scaly=y,smooth=s
 cap-color: brown=n, buff=b, cinnamon=c, gray=g, green=r, pink=p, purple=u, red=e, white=w, yellow=y
 bruises?: bruises=t, no=f
 odor: almond=a, anise=l, creosote=c, fishy=y, foul=f, musty=m, none=n, pungent=p, spicy=s
 gill-attachment: attached=a, descending=d, free=f, notched=n
 gill-spacing: close=c, crowded=w, distant=d
 gill-size: broad=b, narrow=n
 gill-color: black=k, brown=n, buff=b, chocolate=h, gray=g, green=r, orange=o, pink=p, purple=u, red=e, white=w, yellow=y
 stalk-shape: enlarging=e, tapering=t
 stalk-root: bulbous=b, club=c, cup=u, equal=e, rhizomorphs=z, rooted=r, missing=?
 stalk-surface-above-ring: fibrous=f, scaly=y, silky=k, smooth=s
 stalk-surface-below-ring: fibrous=f, scaly=y, silky=k, smooth=s
 stalk-color-above-ring: brown=n, buff=b, cinnamon=c, gray=g, orange=o, pink=p, red=e, white=w, yellow=y
 stalk-color-below-ring: brown=n, buff=b, cinnamon=c, gray=g, orange=o, pink=p, red=e, white=w, yellow=y
 veil-type: partial=p, universal=u
 veil-color: brown=n, orange=o, white=w, yellow=y
 ring-number: none=n, one=o, two=t
 ring-type: cobwebby=c, evanescent=e, flaring=f, large=l, none=n, pendant=p, sheathing=s, zone=z
 spore-print-color: black=k, brown=n, buff=b, chocolate=h, green=r, orange=o, purple=u, white=w, yellow=y
 population: abundant=a, clustered=c, numerous=n, scattered=s, several=v, solitary=y
 habitat: grasses=g, leaves=l, meadows=m, paths=p, urban=u, waste=w, woods=d

sunken is one possible value of the cap-shape feature

A learned decision tree

Root → leaf

```

odor = a: e (400.0)
odor = c: p (192.0)
odor = f: p (2160.0)
odor = l: e (400.0)
odor = m: p (36.0)
odor = n

```

if odor=almond, predict edible

```

    spore-print-color = b: e (48.0)
    spore-print-color = h: e (48.0)
    spore-print-color = k: e (1296.0)
    spore-print-color = n: e (1344.0)
    spore-print-color = o: e (48.0)
    spore-print-color = r: p (72.0)
    spore-print-color = u: e (0.0)
    spore-print-color = w

```

```

        gill-size = b: e (528.0)
        gill-size = n

```

```

            gill-spacing = c: p (32.0)

```

if odor=none $\wedge$
 spore-print-color=white $\wedge$
 gill-size=narrow $\wedge$
 gill-spacing=crowded,
 predict poisonous

```

            gill-spacing = d: e (0.0)

```

```

            gill-spacing = w

```

```

                population = a: e (0.0)
                population = c: p (16.0)
                population = n: e (0.0)
                population = s: e (0.0)
                population = v: e (48.0)
                population = y: e (0.0)

```

```

        spore-print-color = y: e (48.0)

```

```


odor = p: p (256.0)
odor = s: p (576.0)
odor = y: p (576.0)

```


Classification with a learned decision tree



$x = \langle \text{bell}, \text{fibrous}, \text{brown}, \text{false},$
 $\text{foul}, \dots \rangle$



```
odor = a: e (400.0)
odor = c: p (192.0)
odor = f: p (2160.0)
odor = l: e (400.0)
odor = m: p (36.0)
odor = n
  spore-print-color = b: e (48.0)
  spore-print-color = h: e (48.0)
  spore-print-color = k: e (1296.0)
  spore-print-color = n: e (1344.0)
  spore-print-color = o: e (48.0)
  spore-print-color = r: p (72.0)
  spore-print-color = u: e (0.0)
  spore-print-color = w
    gill-size = b: e (528.0)
    gill-size = n
      gill-spacing = c: p (32.0)
      gill-spacing = d: e (0.0)
      gill-spacing = w
        population = a: e (0.0)
        population = c: p (16.0)
        population = n: e (0.0)
        population = s: e (0.0)
        population = v: e (48.0)
        population = y: e (0.0)
      spore-print-color = y: e (48.0)
    odor = p: p (256.0)
    odor = s: p (576.0)
    odor = y: p (576.0)
```

Trained decision tree



$y = p$
 $y = ?$

Unsupervised learning

- in unsupervised learning, we're given a set of instances, without y's



$$\mathbf{x}^{(1)}, \mathbf{x}^{(2)} \dots \mathbf{x}^{(m)}$$

goal: discover interesting regularities/structures/patterns that characterize the instances


- common unsupervised learning tasks
 - *clustering*
 - *anomaly detection*
 - *dimensionality reduction*

Clustering

- given
 - training set of instances $\mathbf{x}^{(1)}, \mathbf{x}^{(2)} \dots \mathbf{x}^{(m)}$
- output
 - model $h \in H$ that divides the training set into clusters such that there is intra-cluster similarity and inter-cluster dissimilarity




Clustering example

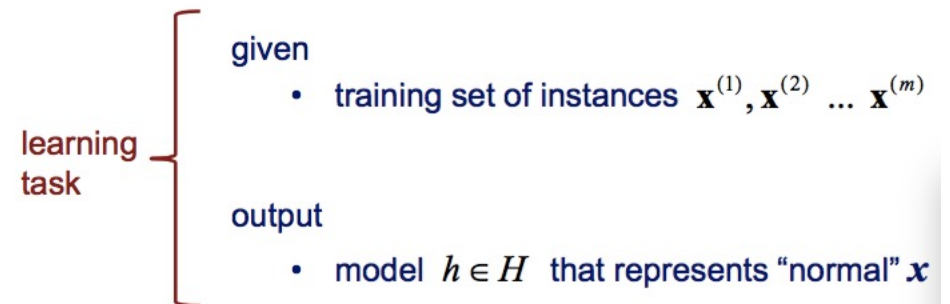
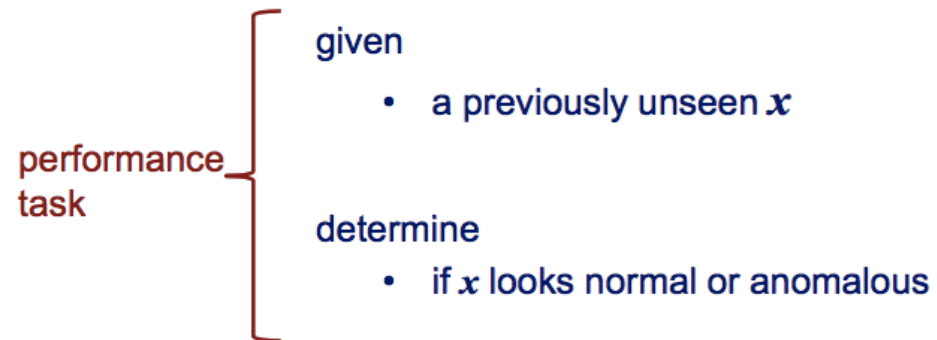


Semi-Supervised Learning


- using labelled as well as unlabelled data to perform certain learning tasks



Anomaly detection



Anomaly detection example




Let's say our model
is represented by:
1979-2000 average,
 ± 2 stddev.

Does the data for
2012 look
anomalous?

Dimensionality reduction

- given
 - training set of instances $\mathbf{x}^{(1)}, \mathbf{x}^{(2)} \dots \mathbf{x}^{(m)}$
- output
 - Model $h \in H$ that represents each $\mathbf{x}$ with a lower-dimension feature vector while still preserving key properties of the data



Dimensionality reduction example



We can represent a face using all of the pixels in a given image



More effective method (for many tasks): represent each face as a linear combination of *eigenfaces*

Dimensionality reduction example

- represent each face as a linear combination of *eigenfaces*

$$\text{face}_1 = \alpha_1^{(1)} \times \text{eigenface}_1 + \alpha_2^{(1)} \times \text{eigenface}_2 + \dots + \alpha_{20}^{(1)} \times \text{eigenface}_{20}$$


$$\mathbf{x}^{(1)} = \langle \alpha_1^{(1)}, \alpha_2^{(1)}, \dots, \alpha_{20}^{(1)} \rangle$$

$$\text{face}_2 = \alpha_1^{(2)} \times \text{eigenface}_1 + \alpha_2^{(2)} \times \text{eigenface}_2 + \dots + \alpha_{20}^{(2)} \times \text{eigenface}_{20}$$


$$\mathbf{x}^{(2)} = \langle \alpha_1^{(2)}, \alpha_2^{(2)}, \dots, \alpha_{20}^{(2)} \rangle$$

- # of features is now 20 instead of # of pixels in images

Other learning tasks



- later in the semester we'll cover other learning tasks that are not strictly supervised or unsupervised
 - *reinforcement learning*
 - *transfer learning*
 - *etc.*



Training Schemes


Batch vs. online learning

- In batch learning, the learner is given the training set as a batch (i.e. all at once)

$$\left(\mathbf{x}^{(1)}, y^{(1)}\right), \left(\mathbf{x}^{(2)}, y^{(2)}\right) \dots \left(\mathbf{x}^{(m)}, y^{(m)}\right)$$



- In online learning, the learner receives instances sequentially, and updates the model after each (for some tasks it might have to classify/make a prediction for each $\mathbf{x}^{(i)}$ before seeing $y^{(i)}$)




Active learning and concept drift

Active learning: cases where the learner can select which instances for training

the target function changes over time
(*concept drift*)

Learning from feedback

- RLHF (Reinforcement Learning from **H**uman Feedback) trains AI by using direct human preferences to guide the model,



How to scale it?

- RLAIFF (Reinforcement Learning from **A**I Feedback) uses an AI model to generate preference labels instead of humans



How to ensure alignment?

Generalization

- The primary objective in supervised learning is to find a model that *generalizes*
 - one that accurately predicts y for previously unseen x

Can I eat this
mushroom that **was**
not in my training set?

